





# Docker for the UseR

*Noam Ross, nyhacker, 2018-07-10*


@noamross



# EcoHealth Alliance




@EcoHealthNYC [ecohealthalliance.org](http://ecohealthalliance.org)




Building technical and community  
infrastructure for R to support  
open, reproducible science

@rOpenSci  
[ropensci.org](http://ropensci.org)




docker?

# Virtual Machine Technologies




# A lighter weight framework for scalability

Virtual Machine



Docker



# Scalability → Ubiquity → Galactic Conquest



# A bare minimum of vocabulary

**image** a package defining a working environment

**container** an instance of an image / a running image

**Dockerfile** a file with a series of commands to build an image

**Docker Hub** a website/repository for sharing Docker images

# A bare minimum of commands

docker pull [*user/repo*]

Get image from a repository

docker build [*directory*]  
-t [*tag*]

Make an image from a Docker file  
and give it a tag/name

docker run [*image*]  
-p [*ports*]  
-v [*volumes*]  
[*command*] -d

Start a container from an image  
opening a set of ports  
linking a set volumes  
run a non-default command, detach

docker stop [*container*]

Stop a running container

# A bare minimum of commands in R

```
library(stevedore) # github.com/richfitz/stevedore
                   `harbor`, and `docker` also options
dc <- docker_client()

dc$image$pull(name = [user/repo])
dc$image$build(context = [directory])
my_container <-
  dc$container$run(image = [image], ports = [ports],
                    volumes = [volumes], cmd = [command],
                    detach = TRUE)
my_container$stop()
```

Now, let's use Docker to ...


Make a fixed working environment

Access a service outside of R

Create an R based service

Send our compute job somewhere else

# Docker as a Fixed Working Environment



09-make-Figure04-viral-trait.R	Generate Figure 4
10-make-ExtendedFigure02-heatmap.R	Generate heat map for Extended Figure 2
11-make-ExtendedTable01-models.R	Generate Extended Table 1 of model summaries
12-make-SuppTable01-predictions.R	Generate supplemental table of observed and predicted values
intermediates/	Holds intermediate fitted model objects when project is run
shapefiles/	Holds large shapefiles downloaded when project is run
packrat/	Holds all R package dependencies
.Rprofile	Configures R to use packrat dependencies

## Reproducing the analysis

The `Makefile` in this repository holds the project workflow. Running `make all` in the directory will re-build the project. `make clean` will remove shapefiles, intermediate data, fit models, and all figures and maps. If this project is opened in RStudio, this can also be accomplished with the "Build All" and "Clean" buttons in the Build tab.

This project uses `packrat` to manage R package dependencies. Running `packrat::restore()` will unpack the versions of packages used in this project. In addition, these packages have the following system requirements: `cairo`, `gdal`, `GEOS`, `libmagick++-`, `jave`, `curl`, `libpng`, `libxml2`, `OpenSSL`, and `pandoc`. All analyses were performed using R 3.3.2 under Ubuntu 14.04. Complete build takes approximately 1 hour with 40 cores and 256GB of memory, or approximately 8 hours on a 2-core Macbook Pro with 16GB of memory.


# Getting people on the same page - within and across teams

nature  
biotechnology

Reproducibility of computational workflows is automated using continuous analysis

Brett K Beaulieu-Jones<sup>1</sup> & Casey S Greene<sup>2</sup>

Beaulieu-Jones & Greene  
(2017) doi:10.1038/nbt.3780





rocker-org

rocker-project.org, [github.com/rocker-org](https://github.com/rocker-org)

Repositories 17

People 2

Projects 0

Search repositories...

Type: All ▾

Language: All ▾

## rocker

R configurations for Docker

Shell ★ 759 ⚡ 205 GPL-2.0 Updated 13 hours ago



### Top languages

Shell Makefile CSS TeX

## rocker-versioned


Run current & prior versions of R using docker

Shell ★ 82 ⚡ 76 GPL-2.0 Updated 5 days ago



### People

2 ▾



cboettig

Carl Boettiger



eddelbuettel

Dirk Eddelbuettel

## shiny

Shell ★ 91 ⚡ 79 Updated 6 days ago

Boettiger + Eddelbuettel (2017) [arXiv:1410.0846](https://arxiv.org/abs/1410.0846)  
Boettiger (2015) doi:10.1145/2723872.2723882

## The versioned stack

image	description	size	metrics	build status
r-ver	Specify R version in docker tag. Builds on <a href="#">debian:stable</a>	221.8MB 8 layers	docker pulls 31k	docker build automated
rstudio	Adds rstudio	336.4MB 17 layers	docker pulls 1M	docker build automated
tidyverse	Adds tidyverse & devtools	625.1MB 18 layers	docker pulls 165k	docker build automated
verse	Adds tex & publishing-related packages	873.3MB 19 layers	docker pulls 47k	docker build automated
geospatial	Adds geospatial libraries	1.3GB 21 layers	docker pulls 23k	docker build automated

## The base stack

r-base	Current R via apt-get with <a href="#">debian:testing</a> & <a href="#">unstable</a> repos	285.3MB 12 layers	docker pulls 2M	docker build automated
r-devel	R-devel added side-by-side onto r-base (using alias <a href="#">RD</a> )	1.2GB 20 layers	docker pulls 5k	docker build automated
drd	lighter r-devel, built not quite daily	646.6MB 16 layers	docker pulls 5k	docker build automated

## Additional images

r-devel-san	as r-devel, but built with compiler sanitizers	1.6GB 21 layers	docker pulls 2k	docker build automated
r-devel-ubsan-clan	Sanitizers, clang c compiler (instead of gcc)	1.8GB 21 layers	docker pulls 789	docker build automated
rstudio:testing	rstudio on debian:testing	1.6GB 21 layers	docker pulls 2k	docker build automated
shiny	shiny-server on r-base	406.1MB 17 layers	docker pulls 186k	docker build automated
r-apt	(R plus CRAN + marutter repo information)	254.5MB 14 layers	docker pulls 3k	docker build automated

**Let's Demo!**

# Reproducibility Helpers: Containerit to Capture Dependencies

## tl;dr

Load the package, do your analysis, and create a Dockerfile.

```
library("containerit")

container <- dockerfile(from = sessionInfo())


cat(as.character(format(container)), sep = "\n")

## FROM rocker/r-ver:3.4.4
## LABEL maintainer="daniel"
## RUN export DEBIAN_FRONTEND=noninteractive; apt-get -y update \
##     && apt-get install -y git-core \
##     libapparmor-dev \
##     libxml2-dev \
##     make \
##     pandoc \
##     pandoc-citeproc
## RUN ["install2.r", "assertthat", "backports", "colorspace", "commonmark", "crayon", "desc", "digest", "e
## RUN ["installGithub.r", "r-lib/pkgdown@5e4825875751c009444c56ce43d06324ec53e910"]
## WORKDIR /payload/
## CMD ["R"]
```



# Reproducibility Helpers: Liftr for Dockerized R Markdown

```
title: "The Missing Example of liftr"
author: "Author Name"
date: "2018-06-26"
output: rmarkdown::html_document
liftr:
  maintainer: "Maintainer Name"
  email: "name@example.com"
  from: "rocker/r-base:latest"
  pandoc: true
  texlive: false
  sysdeps:
    - gfortran
  cran:
    - glmnet
  bioc:
    - Gviz
  remotes:
    - "road2stat/liftr"
```



liftr extends the R Markdown metadata format, introducing additional options for containerizing and rendering reports.

By running `lift()` on the RMD file, liftr parses the metadata fields appeared in the R Markdown document; then generates the Dockerfile.

By running `render_docker()`, liftr will build the Docker image, run the container, and render the R Markdown document.



## Build and launch a repository

GitHub repo or URL

Git branch, tag, or commit



Path to a notebook file (optional)

Copy the URL below and share your Binder with others:




Copy the text below, then paste into your README to show a binder badge: [launch](#) [binder](#)



[Code](#)[Issues 0](#)[Pull requests 0](#)[Projects 0](#)[Wiki](#)[Insights](#)

Branch: master ▾

[noise-phenomena / README.md](#)[Find file](#)[Copy path](#) **cboettig** Update README.md

13a8e02 on May 24

1 contributor

27 lines (16 sloc) | 2.25 KB

[Raw](#)[Blame](#)[History](#)

## noise-phenomena compendium

[launch](#) [binder](#) [build](#) passing [DOI](#) [10.5281/zenodo.1219780](#)

A compendium of code, data, and author's manuscript accompanying the publication:

Carl Boettiger . *From noise to knowledge: how randomness generates novel phenomena and reveals information*. Published in *Ecology Letters*, 22 May 2018 <https://doi.org/10.1111/ele.13085>

# Reservoir: The EHA Data Science Container

A docker image for EHA's modeling and analytics work servers.




*reservoir* is an image built for the modeling and analytics workflow at [EcoHealth Alliance](#). It build on top of the [rocker project geospatial](#) Docker image, and adds:

- A start-up workflow to load user accounts from the host server
- SSH and [mosh](#) access and a suite of command-line tools
- RStudio [preview edition](#)
- Additional system and R libraries for geospatial (GRASS7/rgrass7) and optimization libraries (glpk/Symphony)

See the Dockerfile for details

# Docker as provider of libraries/services



# External library/service examples



Browser rendering services for web scraping (`splashr`, `RSelenium`)




Geospatial Services (`osrm`)



ML Services (`h2o`)

**Let's Demo!**

# Docker as a platform for R services/apps



# Types of R Services



Web Apps



APIs



IDEs

**Let's Demo!**

# rize

417.6MB 22 layers version latest DOI [10.5281/zenodo.322636](https://doi.org/10.5281/zenodo.322636)

A robust method to automatically dockerize your R Shiny Application

## Dockerize Your Shiny App


Within the working directory of a shiny application, run in R:

```
rize::shiny_dockerize()
```

This function takes the following steps to build, test, and view a dockerized shiny application:

1. Make a dependencies file if one does not already exist using the `automagic` R package
2. Create a simple Dockerfile that relies on the `colebrokamp/rize` Docker image
3. Find a valid docker executable with `find_docker_cmd`
4. Build the docker image with `build_docker_app`
5. `view_docker_app` : starts the app in a container and opens in RStudio Viewer or browser


# Docker for deploying compute jobs



# An example with the future package

```
library(future)
library(future.apply)
cl <- makeClusterPSOCK(
  workers = c("10.0.0.40", "10.0.0.43"),
  rscript = c("docker", "run",
             "henrikbengtsson/r-base-future", "Rscript"))
plan(cluster, workers = cl)
result <- future.apply(1:10, a_long_function_run_in_docker)
```

# Docker for deploying compute jobs



# Singularity



Don't need admin rights on computer to use

Less isolation - shares files with your system

For packaging up workflows, scripts, utilities

Targets scientific applications and computing clusters

Runs Docker images, too



Noam Ross

@noamross

Come to my talk at [@nyhackr](#) tomorrow,  
where we'll brave the high seas of [#rstats](#)  
and Docker! [meetup.com/nyhackr/events ...](https://meetup.com/nyhackr/events/)  
7PM at [@ebaynyc](#), 6:30.





**Henrik Bengtsson**  
@henrikbengtsson

Following



Replying to @noamross @nyhackr @ebaynyc

Couldn't resist - here's the [@SingularityApp](#) version 😊



2:54 AM - 10 Jul 2018

# More Resources!


The Dockumentation: [docs.docker.com](https://docs.docker.com)

An rOpenSci Lesson:

[ropenscilabs.github.io/r-docker-tutorial/](https://ropenscilabs.github.io/r-docker-tutorial/)

All links and resources mentioned here will be live tomorrow at:

[github.com/noamross/nyhackr-docker-talk](https://github.com/noamross/nyhackr-docker-talk)



# Thanks!

@noamross

@EcoHealthNYC

@rOpenSci